

## ▼ BERT with Back Translation Augmentation

### Real News and Fake News (~82k total)

Class: Label

Real: 1

Fake: 0

```
import tensorflow as tf

# Get the GPU device name.
device_name = tf.test.gpu_device_name()

# The device name should look like the following:
if device_name == '/device:GPU:0':
    print('Found GPU at: {}'.format(device_name))
else:
    raise SystemError('GPU device not found')

☞ Found GPU at: /device:GPU:0

import torch

# If there's a GPU available...
if torch.cuda.is_available():

    # Tell PyTorch to use the GPU.
    device = torch.device("cuda")

    print('There are %d GPU(s) available.' % torch.cuda.device_count())

    print('We will use the GPU:', torch.cuda.get_device_name(0))

# If not
```

```
# if not...
else:
    print('No GPU available, using the CPU instead.')
    device = torch.device("cpu")
```

```
☞ There are 1 GPU(s) available.
We will use the GPU: Tesla P100-PCIE-16GB
```

```
!pip install transformers
```

## ▼ Back Translation Augmentation

```
# Back Translation
```

```
import pandas as pd
import numpy as np
```

```
import sklearn
from sklearn.model_selection import train_test_split
```

```
file = "combined_{}.csv"
dfs = []
for i in range(3):
    fp = file.format(i+1)
    read = pd.read_csv(fp)
    read = read[['label', 'clean_text']]
    dfs.append(read)
```

```
dfs[2] = dfs[2][:-13000]
data = pd.concat(dfs)
data.tail()
data.reset_index(inplace=True, drop=True)
```

```
data.dropna(inplace=True)
train_data, test_data = train_test_split(data, test_size=0.3)
```

```
print('Train Data:', train_data.shape)
```

```
print(train_data[train_data.label == 1].shape[0], "Real")
print(train_data[train_data.label == 0].shape[0], "Fake")

print('\nTest Data:', test_data.shape)
print(test_data[test_data.label == 1].shape[0], "Real")
print(test_data[test_data.label == 0].shape[0], "Fake")

bt1 = pd.read_csv('Kaggle2_Mixed_bt_clean.csv')[['label', 'clean_text']]
bt2 = pd.read_csv('LIAR_BT_clean.csv')[['label', 'clean_text']]

bt_augmentation = pd.concat([bt1, bt2])
train_data = pd.concat([train_data, bt_augmentation])
train_data.dropna(inplace=True)

print('\n-----After Augmentation ----- \n')
print('Train Data:', train_data.shape)
print(train_data[train_data.label == 1].shape[0], "Real")
print(train_data[train_data.label == 0].shape[0], "Fake")

print('\nTest Data:', test_data.shape)
print(test_data[test_data.label == 1].shape[0], "Real")
print(test_data[test_data.label == 0].shape[0], "Fake")

sentences = train_data.clean_text.values
labels = train_data.label.values

train_data.head(10)
```



```
Train Data: (41621, 2)
22107 Real
19514 Fake
```

```
Test Data: (17838, 2)
9342 Real
8496 Fake
```

```
-----After Augmentation -----
```

```
Train Data: (64435, 2)
33292 Real
31143 Fake
```

```
Test Data: (17838, 2)
9342 Real
8496 Fake
```

	<b>label</b>	<b>clean_text</b>
<b>43534</b>	1	marco rubios economic proposals add trillion ...
<b>15764</b>	0	november eduard popov fort russ translated...
<b>40554</b>	1	president bush eight years added trillion deb...
<b>50229</b>	1	greatest time may losing edge jordan brand b...
<b>23399</b>	0	november pm trump fans may election arent t...
<b>76</b>	0	email get ready cringeworthy story youre going...
<b>46139</b>	1	says jimrenacci consistently voted loopholes e...
<b>7974</b>	0	chart day explosion student loans vs implosion...
<b>15119</b>	0	politics leader islamic revolution ayatollah s...
<b>52126</b>	1	almost us think starting business point though...

```
from transformers import BertTokenizer
```

```
# Load the BERT tokenizer.
```

```

print('Loading BERT tokenizer...')
tokenizer = BertTokenizer.from_pretrained('bert-base-uncased', do_lower_case=True)

# Print the original sentence.
print(' Original: ', sentences[0])

# Print the sentence split into tokens.
print('Tokenized: ', tokenizer.tokenize(sentences[0]))

# Print the sentence mapped to token ids.
print('Token IDs: ', tokenizer.convert_tokens_to_ids(tokenizer.tokenize(sentences[0])))

↳ Loading BERT tokenizer...
   Original: marco rubios economic proposals add trillion federal deficit
   Tokenized: ['marco', 'rub', '##ios', 'economic', 'proposals', 'add', 'trillion', 'federal', 'deficit']
   Token IDs: [8879, 14548, 10735, 3171, 10340, 5587, 23458, 2976, 15074]

# Tokenize all of the sentences and map the tokens to thier word IDs.
input_ids = []

# For every sentence...
for sent in sentences:
    # `encode` will:
    #   (1) Tokenize the sentence.
    #   (2) Prepend the `[CLS]` token to the start.
    #   (3) Append the `[SEP]` token to the end.
    #   (4) Map tokens to their IDs.
    encoded_sent = tokenizer.encode(
        sent,                                     # Sentence to encode.
        add_special_tokens = True, # Add '[CLS]' and '[SEP]'
        max_length = 512 # Truncate all sentences.
        #return_tensors = 'pt',      # Return pytorch tensors.
    )

    # Add the encoded sentence to the list.
    input_ids.append(encoded_sent)

# Print sentence 0, now as a list of IDs.
print('Original: ', sentences[0])

```

```
print('Token IDs:', input_ids[0])
```

```
↳ Original: marco rubios economic proposals add trillion federal deficit  
Token IDs: [101, 8879, 14548, 10735, 3171, 10340, 5587, 23458, 2976, 15074, 102]
```

```
import statistics
```

```
print('Avg sentence length: ', statistics.mean([len(sen) for sen in input_ids]))
```

```
↳ Avg sentence length: 179.84809497943664
```

```
print('Max sentence length: ', max([len(sen) for sen in input_ids]))
```

```
↳ Max sentence length: 512
```

```
import keras
```

```
# We'll borrow the `pad_sequences` utility function to do this.
```

```
from keras.preprocessing.sequence import pad_sequences
```

```
# Set the maximum sequence length.
```

```
# I've chosen 64 somewhat arbitrarily. It's slightly larger than the
```

```
# maximum training sentence length of 47...
```

```
MAX_LEN = 256
```

```
print('\nPadding/truncating all sentences to %d values...' % MAX_LEN)
```

```
print('\nPadding token: "{:}", ID: {:}'.format(tokenizer.pad_token, tokenizer.pad_token_id))
```

```
# Pad our input tokens with value 0.
```

```
# "post" indicates that we want to pad and truncate at the end of the sequence,
```

```
# as opposed to the beginning.
```

```
input_ids = pad_sequences(input_ids, maxlen=MAX_LEN, dtype="long",  
                          value=0, truncating="post", padding="post")
```

```
print('\nDone.')
```

```
↳
```

Padding/truncating all sentences to 256 values...

Padding token: "[PAD]", ID: 0  
Using TensorFlow backend.

Done.

```
# Create attention masks
attention_masks = []

# For each sentence...
for sent in input_ids:

    # Create the attention mask.
    # - If a token ID is 0, then it's padding, set the mask to 0.
    # - If a token ID is > 0, then it's a real token, set the mask to 1.
    att_mask = [int(token_id > 0) for token_id in sent]

    # Store the attention mask for this sentence.
    attention_masks.append(att_mask)

# Use 90% for training and 10% for validation.
train_inputs, validation_inputs, train_labels, validation_labels = train_test_split(input_ids, labels,
                                                                                      random_state=2018, test_size=0.1)

# Do the same for the masks.
train_masks, validation_masks, _, _ = train_test_split(attention_masks, labels,
                                                         random_state=2018, test_size=0.1)

# Convert all inputs and labels into torch tensors, the required datatype
# for our model.
train_inputs = torch.tensor(train_inputs)
validation_inputs = torch.tensor(validation_inputs)

train_labels = torch.tensor(train_labels)
validation_labels = torch.tensor(validation_labels)
```

```

train_masks = torch.tensor(train_masks)
validation_masks = torch.tensor(validation_masks)

from torch.utils.data import TensorDataset, DataLoader, RandomSampler, SequentialSampler

# The DataLoader needs to know our batch size for training, so we specify it
# here.
# For fine-tuning BERT on a specific task, the authors recommend a batch size of
# 16 or 32.

batch_size = 32

# Create the DataLoader for our training set.
train_data = TensorDataset(train_inputs, train_masks, train_labels)
train_sampler = RandomSampler(train_data)
train_dataloader = DataLoader(train_data, sampler=train_sampler, batch_size=batch_size)

# Create the DataLoader for our validation set.
validation_data = TensorDataset(validation_inputs, validation_masks, validation_labels)
validation_sampler = SequentialSampler(validation_data)
validation_dataloader = DataLoader(validation_data, sampler=validation_sampler, batch_size=batch_size)

from transformers import BertForSequenceClassification, AdamW, BertConfig

# Load BertForSequenceClassification, the pretrained BERT model with a single
# linear classification layer on top.
model = BertForSequenceClassification.from_pretrained(
    "bert-base-uncased", # Use the 12-layer BERT model, with an uncased vocab.
    num_labels = 2, # The number of output labels--2 for binary classification.
                    # You can increase this for multi-class tasks.
    output_attentions = False, # Whether the model returns attentions weights.
    output_hidden_states = False, # Whether the model returns all hidden-states.
)

# Tell pytorch to run this model on the GPU.
model.cuda()

```

```

# Note: AdamW is a class from the huggingface library (as opposed to pytorch)

```



```

# I believe the 'W' stands for 'Weight Decay fix"
optimizer = AdamW(model.parameters(),
                    lr = 2e-5, # args.learning_rate - default is 5e-5, our notebook had 2e-5
                    eps = 1e-8 # args.adam_epsilon - default is 1e-8.
                    )

from transformers import get_linear_schedule_with_warmup

# Number of training epochs (authors recommend between 2 and 4)
epochs = 5

# Total number of training steps is number of batches * number of epochs.
total_steps = len(train_dataloader) * epochs

# Create the learning rate scheduler.
scheduler = get_linear_schedule_with_warmup(optimizer,
                                              num_warmup_steps = 0, # Default value in run_glue.py
                                              num_training_steps = total_steps)

# Function to calculate the accuracy of our predictions vs labels
def flat_accuracy(preds, labels):
    pred_flat = np.argmax(preds, axis=1).flatten()
    labels_flat = labels.flatten()
    return np.sum(pred_flat == labels_flat) / len(labels_flat)

import time
import datetime

def format_time(elapsed):
    """
    Takes a time in seconds and returns a string hh:mm:ss
    """
    # Round to the nearest second.
    elapsed_rounded = int(round((elapsed)))

    # Format as hh:mm:ss

```

```

return str(datetime.timedelta(seconds=elapsed_rounded))

torch.cuda.empty_cache()

import random

# This training code is based on the `run_glue.py` script here:
# https://github.com/huggingface/transformers/blob/5bfcd0485ece086ebcbed2d008813037968a9e58/examples/run\_glue.py#I

# Set the seed value all over the place to make this reproducible.
seed_val = 42

random.seed(seed_val)
np.random.seed(seed_val)
torch.manual_seed(seed_val)
torch.cuda.manual_seed_all(seed_val)

# Store the average loss after each epoch so we can plot them.
loss_values = []

# For each epoch...
for epoch_i in range(0, epochs):

    # =====
    #             Training
    # =====

    # Perform one full pass over the training set.

    print("")
    print('==== Epoch {:} / {:} ====='.format(epoch_i + 1, epochs))
    print('Training...')

    # Measure how long the training epoch takes.
    t0 = time.time()

    # Reset the total loss for this epoch.
    total_loss = 0

```

```
total_loss = 0
```

```
# Put the model into training mode. Don't be mislead--the call to
# `train` just changes the *mode*, it doesn't *perform* the training.
# `dropout` and `batchnorm` layers behave differently during training
# vs. test (source: https://stackoverflow.com/questions/51433378/what-does-model-train-do-in-pytorch)
model.train()

# For each batch of training data...
for step, batch in enumerate(train_dataloader):

    # Progress update every 40 batches.
    if step % 40 == 0 and not step == 0:
        # Calculate elapsed time in minutes.
        elapsed = format_time(time.time() - t0)

        # Report progress.
        print(' Batch {:>5,} of {:>5,}. Elapsed: {:.}.'.format(step, len(train_dataloader), elapsed))

    # Unpack this training batch from our dataloader.
    #
    # As we unpack the batch, we'll also copy each tensor to the GPU using the
    # `to` method.
    #
    # `batch` contains three pytorch tensors:
    # [0]: input ids
    # [1]: attention masks
    # [2]: labels
    b_input_ids = batch[0].to(device)
    b_input_mask = batch[1].to(device)
    b_labels = batch[2].to(device)

    # Always clear any previously calculated gradients before performing a
    # backward pass. PyTorch doesn't do this automatically because
    # accumulating the gradients is "convenient while training RNNs".
    # (source: https://stackoverflow.com/questions/48001598/why-do-we-need-to-call-zero-grad-in-pytorch)
    model.zero_grad()

    # Perform a forward pass (evaluate the model on this training batch).
    # This will return the loss (rather than the model output) because we
```

```
# This will return the loss (rather than the model output) because we
# have provided the `labels`.
# The documentation for this `model` function is here:
# https://huggingface.co/transformers/v2.2.0/model\_doc/bert.html#transformers.BertForSequenceClassification
outputs = model(b_input_ids,
                token_type_ids=None,
                attention_mask=b_input_mask,
                labels=b_labels)

# The call to `model` always returns a tuple, so we need to pull the
# loss value out of the tuple.
loss = outputs[0]

# Accumulate the training loss over all of the batches so that we can
# calculate the average loss at the end. `loss` is a Tensor containing a
# single value; the `.item()` function just returns the Python value
# from the tensor.
total_loss += loss.item()

# Perform a backward pass to calculate the gradients.
loss.backward()

# Clip the norm of the gradients to 1.0.
# This is to help prevent the "exploding gradients" problem.
torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)

# Update parameters and take a step using the computed gradient.
# The optimizer dictates the "update rule"--how the parameters are
# modified based on their gradients, the learning rate, etc.
optimizer.step()

# Update the learning rate.
scheduler.step()

# Calculate the average loss over the training data.
avg_train_loss = total_loss / len(train_dataloader)

# Store the loss value for plotting the learning curve.
loss_values.append(avg_train_loss)
```

```

print("")
print("  Average training loss: {0:.2f}".format(avg_train_loss))
print("  Training epoch took: {}".format(format_time(time.time() - t0)))

# =====
#           Validation
# =====
# After the completion of each training epoch, measure our performance on
# our validation set.

print("")
print("Running Validation...")

t0 = time.time()

# Put the model in evaluation mode--the dropout layers behave differently
# during evaluation.
model.eval()

# Tracking variables
eval_loss, eval_accuracy = 0, 0
nb_eval_steps, nb_eval_examples = 0, 0

# Evaluate data for one epoch
for batch in validation_dataloader:

    # Add batch to GPU
    batch = tuple(t.to(device) for t in batch)

    # Unpack the inputs from our dataloader
    b_input_ids, b_input_mask, b_labels = batch

    # Telling the model not to compute or store gradients, saving memory and
    # speeding up validation
    with torch.no_grad():

        # Forward pass, calculate logit predictions.
        # This will return the logits rather than the loss because we have
        # not provided labels.

```

```

# not provided labels.
# token_type_ids is the same as the "segment ids", which
# differentiates sentence 1 and 2 in 2-sentence tasks.
# The documentation for this `model` function is here:
# https://huggingface.co/transformers/v2.2.0/model\_doc/bert.html#transformers.BertForSequenceClassification
outputs = model(b_input_ids,
                 token_type_ids=None,
                 attention_mask=b_input_mask)

# Get the "logits" output by the model. The "logits" are the output
# values prior to applying an activation function like the softmax.
logits = outputs[0]

# Move logits and labels to CPU
logits = logits.detach().cpu().numpy()
label_ids = b_labels.to('cpu').numpy()

# Calculate the accuracy for this batch of test sentences.
tmp_eval_accuracy = flat_accuracy(logits, label_ids)

# Accumulate the total accuracy.
eval_accuracy += tmp_eval_accuracy

# Track the number of batches
nb_eval_steps += 1

# Report the final accuracy for this validation run.
print(" Accuracy: {:.3f}".format(eval_accuracy/nb_eval_steps))
print(" Validation took: {}".format(format_time(time.time() - t0)))

print("")
print("Training complete!")

```



=====  
Epoch 1 / 5  
=====

Training...

Batch	40	of	1,813.	Elapsed:	0:00:30.
Batch	80	of	1,813.	Elapsed:	0:01:01.
Batch	120	of	1,813.	Elapsed:	0:01:31.
Batch	160	of	1,813.	Elapsed:	0:02:01.
Batch	200	of	1,813.	Elapsed:	0:02:31.
Batch	240	of	1,813.	Elapsed:	0:03:02.
Batch	280	of	1,813.	Elapsed:	0:03:32.
Batch	320	of	1,813.	Elapsed:	0:04:02.
Batch	360	of	1,813.	Elapsed:	0:04:32.
Batch	400	of	1,813.	Elapsed:	0:05:03.
Batch	440	of	1,813.	Elapsed:	0:05:33.
Batch	480	of	1,813.	Elapsed:	0:06:03.
Batch	520	of	1,813.	Elapsed:	0:06:33.
Batch	560	of	1,813.	Elapsed:	0:07:04.
Batch	600	of	1,813.	Elapsed:	0:07:34.
Batch	640	of	1,813.	Elapsed:	0:08:04.
Batch	680	of	1,813.	Elapsed:	0:08:35.
Batch	720	of	1,813.	Elapsed:	0:09:05.
Batch	760	of	1,813.	Elapsed:	0:09:35.
Batch	800	of	1,813.	Elapsed:	0:10:06.
Batch	840	of	1,813.	Elapsed:	0:10:36.
Batch	880	of	1,813.	Elapsed:	0:11:06.
Batch	920	of	1,813.	Elapsed:	0:11:36.
Batch	960	of	1,813.	Elapsed:	0:12:07.
Batch	1,000	of	1,813.	Elapsed:	0:12:37.
Batch	1,040	of	1,813.	Elapsed:	0:13:07.
Batch	1,080	of	1,813.	Elapsed:	0:13:38.
Batch	1,120	of	1,813.	Elapsed:	0:14:08.
Batch	1,160	of	1,813.	Elapsed:	0:14:38.
Batch	1,200	of	1,813.	Elapsed:	0:15:09.
Batch	1,240	of	1,813.	Elapsed:	0:15:39.
Batch	1,280	of	1,813.	Elapsed:	0:16:09.
Batch	1,320	of	1,813.	Elapsed:	0:16:39.
Batch	1,360	of	1,813.	Elapsed:	0:17:10.
Batch	1,400	of	1,813.	Elapsed:	0:17:40.
Batch	1,440	of	1,813.	Elapsed:	0:18:10.
Batch	1,480	of	1,813.	Elapsed:	0:18:41.
Batch	1,520	of	1,813.	Elapsed:	0:19:11.
Batch	1,560	of	1,813.	Elapsed:	0:19:41.
Batch	1,600	of	1,813.	Elapsed:	0:20:11.

```
Batch 1,640 of 1,813. Elapsed: 0:20:42.
Batch 1,680 of 1,813. Elapsed: 0:21:12.
Batch 1,720 of 1,813. Elapsed: 0:21:42.
Batch 1,760 of 1,813. Elapsed: 0:22:13.
Batch 1,800 of 1,813. Elapsed: 0:22:43.
```

```
Average training loss: 0.38
Training epoch took: 0:22:52
```

```
Running Validation...
```

```
Accuracy: 0.849
Validation took: 0:00:49
```

```
===== Epoch 2 / 5 =====
```

```
Training...
```

```
Batch 40 of 1,813. Elapsed: 0:00:30.
Batch 80 of 1,813. Elapsed: 0:01:01.
Batch 120 of 1,813. Elapsed: 0:01:31.
Batch 160 of 1,813. Elapsed: 0:02:01.
Batch 200 of 1,813. Elapsed: 0:02:31.
Batch 240 of 1,813. Elapsed: 0:03:02.
Batch 280 of 1,813. Elapsed: 0:03:32.
Batch 320 of 1,813. Elapsed: 0:04:02.
Batch 360 of 1,813. Elapsed: 0:04:33.
Batch 400 of 1,813. Elapsed: 0:05:03.
Batch 440 of 1,813. Elapsed: 0:05:33.
Batch 480 of 1,813. Elapsed: 0:06:03.
Batch 520 of 1,813. Elapsed: 0:06:34.
Batch 560 of 1,813. Elapsed: 0:07:04.
Batch 600 of 1,813. Elapsed: 0:07:34.
Batch 640 of 1,813. Elapsed: 0:08:05.
Batch 680 of 1,813. Elapsed: 0:08:35.
Batch 720 of 1,813. Elapsed: 0:09:05.
Batch 760 of 1,813. Elapsed: 0:09:36.
Batch 800 of 1,813. Elapsed: 0:10:06.
Batch 840 of 1,813. Elapsed: 0:10:36.
Batch 880 of 1,813. Elapsed: 0:11:07.
Batch 920 of 1,813. Elapsed: 0:11:37.
Batch 960 of 1,813. Elapsed: 0:12:07.
Batch 1,000 of 1,813. Elapsed: 0:12:38.
Batch 1,040 of 1,813. Elapsed: 0:13:08.
Batch 1,080 of 1,813. Elapsed: 0:13:38.
Batch 1,120 of 1,813. Elapsed: 0:14:09.
Batch 1,160 of 1,813. Elapsed: 0:14:39.
```



```
Batch 1,200 of 1,813. Elapsed: 0:15:09.
Batch 1,240 of 1,813. Elapsed: 0:15:40.
Batch 1,280 of 1,813. Elapsed: 0:16:10.
Batch 1,320 of 1,813. Elapsed: 0:16:40.
Batch 1,360 of 1,813. Elapsed: 0:17:11.
Batch 1,400 of 1,813. Elapsed: 0:17:41.
Batch 1,440 of 1,813. Elapsed: 0:18:11.
Batch 1,480 of 1,813. Elapsed: 0:18:42.
Batch 1,520 of 1,813. Elapsed: 0:19:12.
Batch 1,560 of 1,813. Elapsed: 0:19:42.
Batch 1,600 of 1,813. Elapsed: 0:20:13.
Batch 1,640 of 1,813. Elapsed: 0:20:43.
Batch 1,680 of 1,813. Elapsed: 0:21:13.
Batch 1,720 of 1,813. Elapsed: 0:21:44.
Batch 1,760 of 1,813. Elapsed: 0:22:14.
Batch 1,800 of 1,813. Elapsed: 0:22:44.
```

Average training loss: 0.25  
Training epoch took: 0:22:53

Running Validation...

Accuracy: 0.874  
Validation took: 0:00:49

=====  
Epoch 3 / 5  
=====

Training...

```
Batch 40 of 1,813. Elapsed: 0:00:30.
Batch 80 of 1,813. Elapsed: 0:01:01.
Batch 120 of 1,813. Elapsed: 0:01:31.
Batch 160 of 1,813. Elapsed: 0:02:01.
Batch 200 of 1,813. Elapsed: 0:02:31.
Batch 240 of 1,813. Elapsed: 0:03:02.
Batch 280 of 1,813. Elapsed: 0:03:32.
Batch 320 of 1,813. Elapsed: 0:04:02.
Batch 360 of 1,813. Elapsed: 0:04:33.
Batch 400 of 1,813. Elapsed: 0:05:03.
Batch 440 of 1,813. Elapsed: 0:05:33.
Batch 480 of 1,813. Elapsed: 0:06:03.
Batch 520 of 1,813. Elapsed: 0:06:34.
Batch 560 of 1,813. Elapsed: 0:07:04.
Batch 600 of 1,813. Elapsed: 0:07:34.
Batch 640 of 1,813. Elapsed: 0:08:04.
Batch 680 of 1,813. Elapsed: 0:08:35.
Batch 720 of 1,813. Elapsed: 0:09:05.
```

```

Batch 720 of 1,813. Elapsed: 0:09:05.
Batch 760 of 1,813. Elapsed: 0:09:35.
Batch 800 of 1,813. Elapsed: 0:10:05.
Batch 840 of 1,813. Elapsed: 0:10:36.
Batch 880 of 1,813. Elapsed: 0:11:06.
Batch 920 of 1,813. Elapsed: 0:11:36.
Batch 960 of 1,813. Elapsed: 0:12:07.
Batch 1,000 of 1,813. Elapsed: 0:12:37.
Batch 1,040 of 1,813. Elapsed: 0:13:07.
Batch 1,080 of 1,813. Elapsed: 0:13:37.
Batch 1,120 of 1,813. Elapsed: 0:14:08.
Batch 1,160 of 1,813. Elapsed: 0:14:38.
Batch 1,200 of 1,813. Elapsed: 0:15:08.
Batch 1,240 of 1,813. Elapsed: 0:15:39.
Batch 1,280 of 1,813. Elapsed: 0:16:09.
Batch 1,320 of 1,813. Elapsed: 0:16:39.
Batch 1,360 of 1,813. Elapsed: 0:17:09.
Batch 1,400 of 1,813. Elapsed: 0:17:40.
Batch 1,440 of 1,813. Elapsed: 0:18:10.
Batch 1,480 of 1,813. Elapsed: 0:18:40.
Batch 1,520 of 1,813. Elapsed: 0:19:11.
Batch 1,560 of 1,813. Elapsed: 0:19:41.
Batch 1,600 of 1,813. Elapsed: 0:20:11.
Batch 1,640 of 1,813. Elapsed: 0:20:42.
Batch 1,680 of 1,813. Elapsed: 0:21:12.
Batch 1,720 of 1,813. Elapsed: 0:21:42.
Batch 1,760 of 1,813. Elapsed: 0:22:12.
Batch 1,800 of 1,813. Elapsed: 0:22:43.

```

Average training loss: 0.17  
 Training epoch took: 0:22:52

Running Validation...

Accuracy: 0.894  
 Validation took: 0:00:49

==== Epoch 4 / 5 =====

Training...

```

Batch 40 of 1,813. Elapsed: 0:00:30.
Batch 80 of 1,813. Elapsed: 0:01:01.
Batch 120 of 1,813. Elapsed: 0:01:31.
Batch 160 of 1,813. Elapsed: 0:02:01.
Batch 200 of 1,813. Elapsed: 0:02:32.
Batch 240 of 1,813. Elapsed: 0:03:02.

```

Batch	280	of	1,813.	Elapsed:	0:03:32.
Batch	320	of	1,813.	Elapsed:	0:04:03.
Batch	360	of	1,813.	Elapsed:	0:04:33.
Batch	400	of	1,813.	Elapsed:	0:05:03.
Batch	440	of	1,813.	Elapsed:	0:05:33.
Batch	480	of	1,813.	Elapsed:	0:06:04.
Batch	520	of	1,813.	Elapsed:	0:06:34.
Batch	560	of	1,813.	Elapsed:	0:07:04.
Batch	600	of	1,813.	Elapsed:	0:07:35.
Batch	640	of	1,813.	Elapsed:	0:08:05.
Batch	680	of	1,813.	Elapsed:	0:08:35.
Batch	720	of	1,813.	Elapsed:	0:09:06.
Batch	760	of	1,813.	Elapsed:	0:09:36.
Batch	800	of	1,813.	Elapsed:	0:10:06.
Batch	840	of	1,813.	Elapsed:	0:10:37.
Batch	880	of	1,813.	Elapsed:	0:11:07.
Batch	920	of	1,813.	Elapsed:	0:11:37.
Batch	960	of	1,813.	Elapsed:	0:12:07.
Batch	1,000	of	1,813.	Elapsed:	0:12:38.
Batch	1,040	of	1,813.	Elapsed:	0:13:08.
Batch	1,080	of	1,813.	Elapsed:	0:13:38.
Batch	1,120	of	1,813.	Elapsed:	0:14:09.
Batch	1,160	of	1,813.	Elapsed:	0:14:39.
Batch	1,200	of	1,813.	Elapsed:	0:15:09.
Batch	1,240	of	1,813.	Elapsed:	0:15:40.
Batch	1,280	of	1,813.	Elapsed:	0:16:10.
Batch	1,320	of	1,813.	Elapsed:	0:16:40.
Batch	1,360	of	1,813.	Elapsed:	0:17:11.
Batch	1,400	of	1,813.	Elapsed:	0:17:41.
Batch	1,440	of	1,813.	Elapsed:	0:18:11.
Batch	1,480	of	1,813.	Elapsed:	0:18:42.
Batch	1,520	of	1,813.	Elapsed:	0:19:12.
Batch	1,560	of	1,813.	Elapsed:	0:19:42.
Batch	1,600	of	1,813.	Elapsed:	0:20:13.
Batch	1,640	of	1,813.	Elapsed:	0:20:43.
Batch	1,680	of	1,813.	Elapsed:	0:21:13.
Batch	1,720	of	1,813.	Elapsed:	0:21:44.
Batch	1,760	of	1,813.	Elapsed:	0:22:14.
Batch	1,800	of	1,813.	Elapsed:	0:22:44.

Average training loss: 0.11

Training epoch took: 0:22:53

Running Validation...

Accuracy: 0.905

Validation took: 0:00:49

=====  
Epoch 5 / 5  
=====  
Training...

Batch	40	of	1,813.	Elapsed:	0:00:30.
Batch	80	of	1,813.	Elapsed:	0:01:01.
Batch	120	of	1,813.	Elapsed:	0:01:31.
Batch	160	of	1,813.	Elapsed:	0:02:01.
Batch	200	of	1,813.	Elapsed:	0:02:32.
Batch	240	of	1,813.	Elapsed:	0:03:02.
Batch	280	of	1,813.	Elapsed:	0:03:32.
Batch	320	of	1,813.	Elapsed:	0:04:03.
Batch	360	of	1,813.	Elapsed:	0:04:33.
Batch	400	of	1,813.	Elapsed:	0:05:03.
Batch	440	of	1,813.	Elapsed:	0:05:34.
Batch	480	of	1,813.	Elapsed:	0:06:04.
Batch	520	of	1,813.	Elapsed:	0:06:34.
Batch	560	of	1,813.	Elapsed:	0:07:04.
Batch	600	of	1,813.	Elapsed:	0:07:35.
Batch	640	of	1,813.	Elapsed:	0:08:05.
Batch	680	of	1,813.	Elapsed:	0:08:35.
Batch	720	of	1,813.	Elapsed:	0:09:06.
Batch	760	of	1,813.	Elapsed:	0:09:36.
Batch	800	of	1,813.	Elapsed:	0:10:06.
Batch	840	of	1,813.	Elapsed:	0:10:37.
Batch	880	of	1,813.	Elapsed:	0:11:07.
Batch	920	of	1,813.	Elapsed:	0:11:37.
Batch	960	of	1,813.	Elapsed:	0:12:08.
Batch	1,000	of	1,813.	Elapsed:	0:12:38.
Batch	1,040	of	1,813.	Elapsed:	0:13:08.
Batch	1,080	of	1,813.	Elapsed:	0:13:39.
Batch	1,120	of	1,813.	Elapsed:	0:14:09.
Batch	1,160	of	1,813.	Elapsed:	0:14:39.
Batch	1,200	of	1,813.	Elapsed:	0:15:10.
Batch	1,240	of	1,813.	Elapsed:	0:15:40.
Batch	1,280	of	1,813.	Elapsed:	0:16:10.
Batch	1,320	of	1,813.	Elapsed:	0:16:41.
Batch	1,360	of	1,813.	Elapsed:	0:17:11.
Batch	1,400	of	1,813.	Elapsed:	0:17:41.
Batch	1,440	of	1,813.	Elapsed:	0:18:11.
Batch	1,480	of	1,813.	Elapsed:	0:18:42.
Batch	1,520	of	1,813.	Elapsed:	0:19:12.
Batch	1,560	of	1,813.	Elapsed:	0:19:42.

```

Batch 1,560 of 1,813. Elapsed: 0:19:42.
Batch 1,600 of 1,813. Elapsed: 0:20:13.
Batch 1,640 of 1,813. Elapsed: 0:20:43.
Batch 1,680 of 1,813. Elapsed: 0:21:13.
Batch 1,720 of 1,813. Elapsed: 0:21:44.
Batch 1,760 of 1,813. Elapsed: 0:22:14.
Batch 1,800 of 1,813. Elapsed: 0:22:44.

```

```

Average training loss: 0.07
Training epoch took: 0:22:54

```

```

Running Validation...
Accuracy: 0.909
Validation took: 0:00:49

```

```

Training complete!

```

```

# Report the number of sentences.
print('Number of test sentences: {:,}\n'.format(test_data.shape[0]))

# Create sentence and label lists
sentences = test_data.clean_text.values
labels = test_data.label.values

# Tokenize all of the sentences and map the tokens to their word IDs.
input_ids = []

# For every sentence...
for sent in sentences:
    # `encode` will:
    # (1) Tokenize the sentence.
    # (2) Prepend the `[CLS]` token to the start.
    # (3) Append the `[SEP]` token to the end.
    # (4) Map tokens to their IDs.
    encoded_sent = tokenizer.encode(
        sent,                                     # Sentence to encode.
        max_length = 512,
        add_special_tokens = True, # Add '[CLS]' and '[SEP]'
    )

```

```

input_ids.append(encoded_sent)

# Pad our input tokens
input_ids = pad_sequences(input_ids, maxlen=MAX_LEN,
                           dtype="long", truncating="post", padding="post")

# Create attention masks
attention_masks = []

# Create a mask of 1s for each token followed by 0s for padding
for seq in input_ids:
    seq_mask = [float(i>0) for i in seq]
    attention_masks.append(seq_mask)

# Convert to tensors.
prediction_inputs = torch.tensor(input_ids)
prediction_masks = torch.tensor(attention_masks)
prediction_labels = torch.tensor(labels)

# Set the batch size.
batch_size = 32

# Create the DataLoader.
prediction_data = TensorDataset(prediction_inputs, prediction_masks, prediction_labels)
prediction_sampler = SequentialSampler(prediction_data)
prediction_dataloader = DataLoader(prediction_data, sampler=prediction_sampler, batch_size=batch_size)

☞ Number of test sentences: 17,838

# Prediction on test set

print('Predicting labels for {:,} test sentences...'.format(len(prediction_inputs)))

# Put model in evaluation mode
model.eval()

# Tracking variables
predictions, true_labels = [], []

```

```

predictions, true_labels = [], []
test_loss, test_accuracy = 0, 0
nb_test_steps, nb_test_examples = 0, 0

# Predict
for batch in prediction_dataloader:
    # Add batch to GPU
    batch = tuple(t.to(device) for t in batch)

    # Unpack the inputs from our dataloader
    b_input_ids, b_input_mask, b_labels = batch

    # Telling the model not to compute or store gradients, saving memory and
    # speeding up prediction
    with torch.no_grad():
        # Forward pass, calculate logit predictions
        outputs = model(b_input_ids, token_type_ids=None,
                        attention_mask=b_input_mask)

    logits = outputs[0]

    # Move logits and labels to CPU
    logits = logits.detach().cpu().numpy()
    label_ids = b_labels.to('cpu').numpy()

    # # Store predictions and true labels
    # predictions.append(logits)
    # true_labels.append(label_ids)

    # Calculate the accuracy for this batch of test sentences.
    tmp_test_accuracy = flat_accuracy(logits, label_ids)

    # Accumulate the total accuracy.
    test_accuracy += tmp_test_accuracy

    # Track the number of batches
    nb_test_steps += 1

# Report the final accuracy for this validation run.
print("Testing Accuracy: {:.3f}".format(test_accuracy/nb_test_steps))

```

```
print('Testing Accuracy: ', format(testing_accuracy, '%.3f'), '\n\nDONE.',  
      print('    DONE.'))
```

☞ Predicting labels for 17,838 test sentences...  
Testing Accuracy: 0.946  
DONE.